

# Measuring directory size using stat system call

Kaynak: <http://web.eecs.utk.edu/~plank/plank/classes/cs360/>

# Recursive directory traversal

- This lecture covers the writing of a command **prsize**
- What **prsize** does is return the number of bytes taken up by all files reachable from the current directory (excluding soft links).
- Prsize illustrates using **opendir/readdir/closedir**, **stat**, recursion, building path names, and finding hard links.

# prsize1

Open working  
(current)  
directory

Get inode  
information  
of selected  
file

Get file size  
from the  
inode  
information.

```
File Edit View Search Tools Documents Help
Open Save Undo
prsize1.c x
#include <stdio.h>
#include <stdlib.h>
#include <dirent.h>
#include <sys/stat.h>
main()
{
    DIR *d;
    struct dirent *de;
    struct stat buf;
    int exists;
    int total_size;

    d = opendir(".");
    if (d == NULL) {
        perror("prsize");
        exit(1);
    }
    total_size = 0;
    for (de = readdir(d); de != NULL; de = readdir(d)) {
        exists = stat(de->d_name, &buf);
        if (exists < 0) {
            fprintf(stderr, "Couldn't stat %s\n", de->d_name);
        } else {
            total_size += buf.st_size;
        }
    }
    closedir(d);
    printf("%d\n", total_size);
}
```

# prsize1

We temporarily added current directory to PATH so that we can able to use **prsize** programs in other directories.

```
File Edit View Search Terminal Help
abc:3_Prsize$ ./prsize1
367799
abc:3_Prsize$ prsize1
prsize1: command not found
abc:3_Prsize$ PATH=$PATH:$(pwd)
abc:3_Prsize$ prsize1
367799
abc:3_Prsize$ mkdir test1
abc:3_Prsize$ cd test1
abc:test1$ prsize1
8192
abc:test1$
```

# prsize2 function definition

The program that lists the file sizes has been defined as a function. It will help us for recursive directory traversal.

```
prsize2.c x
int get_size(char *fn)
{
    DIR *d;
    struct dirent *de;
    struct stat buf;
    int exists;
    int total_size;

    d = opendir(fn);
    if (d == NULL) {
        perror("prsize");
        exit(1);
    }

    total_size = 0;

    for (de = readdir(d); de != NULL; de = readdir(d)) {
        exists = stat(de->d_name, &buf);
        if (exists < 0) {
            fprintf(stderr, "Couldn't stat %s\n", de->d_name);
        } else {
            total_size += buf.st_size;
        }
    }

    closedir(d);
    return total_size;
}

main()
{
    printf("%d\n", get_size("."));
}
```

# prsize3 : recursive implementation

- Whenever we encounter a directory, we want to find out the size of everything in that directory, so we will call `get_size()` recursively on that directory.
- `S_ISDIR()` checks if a file directory or not.
- If it is a directory, all the files it contains should be measured.
- But this results in the error you see below.

File Edit View Search Terminal Help

test:prsize3

prsize: Too many open files

test:█

```
prsize.c
int get_size(char *fn)
{
    DIR *d;
    struct dirent *de;
    struct stat buf;
    int exists;
    int total_size;

    d = opendir(fn);
    if (d == NULL) {
        perror("prsize");
        exit(1);
    }

    total_size = 0;

    for (de = readdir(d); de != NULL; de = readdir(d)) {
        exists = stat(de->d_name, &buf);
        if (exists < 0) {
            fprintf(stderr, "Couldn't stat %s\n", de->d_name);
        } else {
            total_size += buf.st_size;

            /* Make the recursive call if the file is a directory */
            if (S_ISDIR(buf.st_mode)) {
                total_size += get_size(de->d_name);
            }
        }
        closedir(d);
    }
    return total_size;
}

main()
{
    printf("%d\n", get_size("."));
}
```

## prsize3a: analysis of error

- I put a print statement into [prsize3a.c](#) to see when it's making the recursive calls:
- When the program is run (`./prsize3a`) repeatedly calls the `“.”` (current) directory.
- In other words it calls itself infinitely.
- This goes into an infinite loop until you run out of open file descriptors at which point `opendir()` fails.

File Edit View Search Tools Documents Help

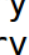
```
total_size = 0;

for (de = readdir(d); de != NULL; de = readdir(d)) {
    exists = stat(de->d_name, &buf);
    if (exists < 0) {
        fprintf(stderr, "Couldn't stat %s\n", de->d_name);
    } else {
        total_size += buf.st_size;
    }
    /* Make the recursive call if the file is a directory */
    if (S_ISDIR(buf.st_mode)) {
        printf("Making recursive call on directory %s\n", de->d_name);
        total_size += get_size(de->d_name);
    }
}

closedir(d);
return total_size;
```

```
main / \
```

```
Making recursive call on directory .  
Making recursive call on directory .  
Making recursive call on directory .  
Making recursive call on directory .  
Making recursive call on directory .  
Making recursive call on directory .  
prsize: Too many open files
```



C ▾ Tab Width: 8 ▾ Ln 38, Col 42 INS

# prsize4 : excluding "." and ".." directories

- If we exclude "." and ".." inside recursive loop program doesn't go into infinite loop.

```
File Edit View Search Tools Documents Help
Open Save Undo
prsize4.c x

for (de = readdir(d); de != NULL; de = readdir(d)) {
    exists = stat(de->d_name, &buf);
    if (exists < 0) {
        fprintf(stderr, "Couldn't stat %s\n", de->d_name);
    } else {
        total_size += buf.st_size;
    }
    /* Make the recursive call if the file is a directory and is not
     * . or .. */
    if (S_ISDIR(buf.st_mode) && strcmp(de->d_name, ".") != 0 &&
        strcmp(de->d_name, "..") != 0) {
        total_size += get_size(de->d_name);
    }
}
closedir(d);
return total_size;

C Tab Width: 8 Ln 3

File Edit View Search Terminal Help
test4:ls -la
total 20
drwxr-xr-x 2 root root 4096 Mar 28 21:44 .
drwx----- 6 bilg bilg 4096 Mar 28 21:43 ..
-rw-r--r-- 1 root root 9 Mar 28 21:44 f1.txt
-rw-r--r-- 1 root root 18 Mar 28 21:44 f2.txt
-rw-r--r-- 1 root root 51 Mar 28 21:44 f3.txt
test4:prsize4
8270
test4:
```



# prsize4 : path problem

- We solved infinite loop problem but there remains another problem with file paths.
- It is necessary to write full path starting from «.» (working directory) before doing a recursive call.

```
test4:touch altklasor/deneme.txt
test4:ls
altklasor  f1.txt  f2.txt  f3.txt
test4:prsize4
Couldn't stat deneme.txt
prsize: No such file or directory
test4:cd altklasor
altklasor:ls
deneme.txt
altklasor:
```

# prsize5: solution to path problem

256 character file  
name + 2 character (/ and null) + file  
path(strlen(fn))

Print the file name to *s*  
variable together with  
file path.

```
File Edit View Search Tools Documents Help
prsize5.c x
char *s;
d = opendir(fn);
if (d == NULL) {
    perror("prsize");
    exit(1);
}
s = (char *) malloc(sizeof(char)*(strlen(fn)+258));

for (de = readdir(d); de != NULL; de = readdir(d)) {
    /* Look for fn/de->d name */
    sprintf(s, "%s/%s", fn, de->d name);
    exists = stat(s, &buf);
    if (exists < 0) {
        fprintf(stderr, "Couldn't stat %s\n", s);
    } else {
        total_size += buf.st_size;
    }
    if (S_ISDIR(buf.st_mode) && strcmp(de->d_name, ".") != 0 &&
        strcmp(de->d_name, "..") != 0) {
        total_size += get_size(s);
    }
}
closedir(d);
```

Saving file '/home/bilg/Documents/h... C Tab Width: 8 Ln 20, Col 11 INS

# prsize5: repeating files problem

- Prsize5 computed the size of directory as 20558.
- However, the files with some inode numbers are taken into account two times.
- What we need is for **prsize** to be able to recognize hard links, and only count them once.
- How do you recognize whether two files are links to the same disk file?
- You use the inode number. This is held in **buf.st\_ino**.

```
test4:prsize5
20558
test4:ls -lai
total 24
1840938 drwxr-xr-x 3 root root 4096 Mar 28 22:06 .
1840034 drwx----- 6 bilg bilg 4096 Mar 28 22:18 ..
1839719 drwxr-xr-x 2 root root 4096 Mar 28 22:06 altklasor
1841007 -rw-r--r-- 1 root root    9 Mar 28 21:44 f1.txt
1841008 -rw-r--r-- 1 root root   18 Mar 28 21:44 f2.txt
1841009 -rw-r--r-- 1 root root   51 Mar 28 21:44 f3.txt
test4:ls -lai altklasor
total 8
1839719 drwxr-xr-x 2 root root 4096 Mar 28 22:06 .
1840938 drwxr-xr-x 3 root root 4096 Mar 28 22:06 ..
1840932 -rw-r--r-- 1 root root    0 Mar 28 22:06 deneme.txt
test4:
```

# prsize6: solution to repeating files problem

- In order to prevent re-use of the same folders, the inode number of each file included in the calculation can be kept in a BST and used for control purposes.
- The red-black tree in the libfdr library is used to keep the inode list.
- It prevents previously used inode for getting the size of a file to be added to the tree again

```
File Edit View Search Tools Documents Help
Open Save Undo
prsize6.c x
#include <stdio.h>
#include <stdlib.h>
#include <string.h>
#include <dirent.h>
#include <sys/stat.h>
#include "jrb.h"

int get_size(char *fn, JRB inodes)
{
    DIR *d;
    struct dirent *de;
    struct stat buf;
    int exists;
    int total_size;
    char *s;

    d = opendir(fn);
    if (d == NULL) {
        perror("prsize");
        exit(1);
    }

    total_size = 0;
    s = (char *) malloc(sizeof(char)*(strlen(fn)+258));

    for (de = readdir(d); de != NULL; de = readdir(d)) {
        /* Look for fn/de->d_name inodes*/
        sprintf(s, "%s/%s", fn, de->d_name);
        exists = stat(s, &buf);
        if (exists == 0) {
```

C Tab Width: 8 Ln 12, Col 1 INS

# prsize6: solution to repeating files problem

```
prsize6.c x
s = (char *) malloc(sizeof(char)*(strlen(fn)+258));

for (de = readdir(d); de != NULL; de = readdir(d)) {
    /* Look for fn/de->d_name inodes */
    sprintf(s, "%s/%s", fn, de->d_name);
    exists = stat(s, &buf);
    if (exists < 0) {
        fprintf(stderr, "Couldn't stat %s\n", s);
    } else {
        if (jrb_find_int(inodes, buf.st_ino) == NULL) {
            total_size += buf.st_size;
            jrb_insert_int(inodes, buf.st_ino, JNULL);
        }
        if (S_ISDIR(buf.st_mode) && strcmp(de->d_name, ".") != 0 &
            strcmp(de->d_name, "..") != 0) {
            total_size += get_size(s, inodes);
        }
    }
    closedir(d);
    free(s);
    return total_size;
}

main()
{
    JRB inodes;
    inodes = make_jrb();
    printf("%d\n", get_size(".", inodes));
}
```

- Search for the inode in tree. Insert if not in the tree.
- Inode is added as key to prevent repetition.
- We don't use value (value=JNULL)

# prsize7.c: soft (symbolic) link problem

We used the lstat ()  
system call for soft  
links.

```
prsize7.c x
S = (cnar *) malloc(sizeof(cnar)*(strlen(Tn)+258));

for (de = readdir(d); de != NULL; de = readdir(d)) {
    /* Look for fn/de->d_name */
    sprintf(s, "%s/%s", fn, de->d_name);
    exists = lstat(s, &buf);
    if (exists < 0) {
        fprintf(stderr, "Couldn't stat %s\n", s);
    } else {
        if (jrb_find_int(inodes, buf.st_ino) == NULL) {
```

The result  
using stat()

The result  
using lstat()

```
File Edit View Search Terminal Help
test4:ln -s f1.txt softf1
test4:ls
altklasor f1.txt f2.txt f3.txt softf1
test4:prsize6
12366
test4:prsize7
12372
test4:ls -li
total 16
1839719 drwxr-xr-x 2 root root 4096 Mar 28 22:06 altklasor
1841007 -rw-r--r-- 1 root root 9 Mar 28 21:44 f1.txt
1841008 -rw-r--r-- 1 root root 18 Mar 28 21:44 f2.txt
1841009 -rw-r--r-- 1 root root 51 Mar 28 21:44 f3.txt
1841073 lrwxrwxrwx 1 root root 6 Mar 29 00:36 softf1 -> f1.txt
test4:
```

# prsize7a: printing visited paths

File Edit View Search Terminal Help

```
abc:1_Prsize$ ./directory 10
```

```
abc:1_Prsize$ ./prsize7a
```

```
Testing .
```

```
Testing ./d10
```

```
Testing ./d10/d9
```

```
Testing ./d10/d9/d8
```

```
Testing ./d10/d9/d8/d7
```

```
Testing ./d10/d9/d8/d7/d6
```

```
Testing ./d10/d9/d8/d7/d6/d5
```

```
Testing ./d10/d9/d8/d7/d6/d5/d4
```

```
Testing ./d10/d9/d8/d7/d6/d5/d4/d3
```

```
Testing ./d10/d9/d8/d7/d6/d5/d4/d3/d2
```

```
Testing ./d10/d9/d8/d7/d6/d5/d4/d3/d2/d1
```

```
Testing ./test1
```

```
413002
```

- Program çalıştırıldığı klasörden (".") başlayıp tüm alt klasörlere aynı işlemi uyguluyor.
- Bir alt klasöre gidilirken önceki klasör açık kalmasıdır.

```
#!/bin/sh
x=$1
while [ $x -gt 0 ]
do
    mkdir d$x
    cd d$x
    echo "test" > f$x
    x=`expr $x - 1`
done
```

sh Tab Width: 8 Ln 1, Col 8 INS



# prsize8: closing a directory before visiting another

- To reduce the number of files opened from a process, the directory can be closed before going to subfolders.
- In this implementation, subdirectories in the visited directory are added to the list.
- Then directory is closed and the subdirectories are visited.
- The same is true for every visited directory.

```
#include <stdio.h>
#include <stdlib.h>
#include <string.h>
#include <dirent.h>
#include <sys/stat.h>
#include "jrb.h"
#include "dllist.h"

int get_size(char *fn, JRB inodes)
{
    DIR *d;
    struct dirent *de;
    struct stat buf;
    int exists;
    int total_size;
    char *s;
    Dllist directories, tmp;

    d = opendir(fn);
    if (d == NULL) {
        perror("prsize");
        exit(1);
    }

    total_size = 0;
    s = (char *) malloc(sizeof(char)*(strlen(fn)+258));

    directories = new_dllist();
    for (de = readdir(d); de != NULL; de = readdir(d)) {
        /* Look for fn/de->d name */
```



# prsize8: closing a directory before visiting another

Add to the list  
if there are  
subdirectories

Close directory  
before going to  
another one

```
directories = new_dlist();
for (de = readdir(d); de != NULL; de = readdir(d)) {
    /* Look for fn/de->d_name */
    sprintf(s, "%s/%s", fn, de->d_name);
    exists = lstat(s, &buf);
    if (exists < 0) {
        fprintf(stderr, "Couldn't stat %s\n", s);
    } else {
        if (jrb_find_int(inodes, buf.st_ino) == NULL) {
            total_size += buf.st_size;
            jrb_insert_int(inodes, buf.st_ino, JNULL);
        }
    }
    if (S_ISDIR(buf.st_mode) && strcmp(de->d_name, ".") != 0 &&
        strcmp(de->d_name, "..") != 0) {
        dll_append(directories, new_jval_s(strdup(s)));
    }
}
closedir(d);
dll_traverse(tmp, directories) {
    total_size += get_size(tmp->val.s, inodes);
    /* This keeps the program from overgrowing its memory */
    free(tmp->val.s);
}

/* As does this */
free_dlist(directories);
free(s);
return total_size;
}

main()
{
    JRB inodes;
    inodes = make_jrb();
    printf("%d\n", get_size(".", inodes));
}
```

Clean after  
getting size